

Module 2: Linear Regression and Gradient Descent

Overview

This module introduces the fundamental algorithm for learning model parameters: gradient descent. Using linear regression as our example, we explore how to define a loss function, compute gradients, and iteratively update parameters to minimize error. These concepts form the basis for training nearly all modern machine learning models.

1. Linear Regression

The Model

Linear regression predicts a continuous output as a weighted sum of inputs:

$$y = w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b$$

Or in vector notation:

$$\mathbf{y} = \mathbf{w} \cdot \mathbf{x} + b$$

Where: - \mathbf{x} : Input features (what we observe) - \mathbf{w} : Weights (learned parameters) - b : Bias/intercept (learned parameter) - y : Predicted output

Why Linear?

Linear models are:

- Simple to understand and interpret
- Fast to train
- Often surprisingly effective
- Foundation for more complex models

Limitations

Linear regression assumes a linear relationship between inputs and outputs. If the true relationship is non-linear (curved), linear regression will underfit.

2. Loss Functions

Purpose of Loss

A loss function measures how wrong our predictions are. We need this to know:

- How good/bad the current model is
- Whether we're improving during training
- Which direction to adjust parameters

Mean Squared Error (MSE)

The most common loss for regression:

$$\text{MSE} = (1/n) \times \sum (y_{\text{pred}} - y_{\text{true}})^2$$

Properties: - Always non-negative (zero is perfect) - Squaring makes all errors positive - Large errors are penalized more than small ones - Differentiable (important for gradient descent)

The Loss Landscape

Imagine the loss as a surface over the space of all possible parameter values. Training is about finding the lowest point on this surface. For linear regression with MSE, this surface is a convex bowl—there's exactly one minimum.

3. Gradient Descent

The Core Idea

Gradient descent is an optimization algorithm that finds parameters minimizing the loss by: 1. Start at a random point in parameter space 2. Compute the gradient (slope) of the loss 3. Take a step in the direction that decreases loss 4. Repeat until convergence

The Gradient

The gradient is a vector of partial derivatives—it points in the direction of steepest increase. To minimize loss, we move in the **opposite** direction.

For MSE loss:

$$\begin{aligned} L/w &= (2/n) \times \sum (y_{\text{pred}} - y_{\text{true}}) \times x \\ L/b &= (2/n) \times \sum (y_{\text{pred}} - y_{\text{true}}) \end{aligned}$$

The Update Rule

$$\begin{aligned} w_{\text{new}} &= w_{\text{old}} - \eta \times (L/w) \\ b_{\text{new}} &= b_{\text{old}} - \eta \times (L/b) \end{aligned}$$

Where η (eta) is the learning rate.

4. Learning Rate

What It Controls

The learning rate determines the step size in each update: - **Too high:** Steps are too big, may overshoot the minimum and diverge - **Too low:** Steps are too small, convergence is very slow - **Just right:** Smooth, efficient convergence

Symptoms of Bad Learning Rates

Too High: - Loss increases instead of decreases - Loss oscillates wildly - Loss becomes NaN or infinity

Too Low: - Loss decreases very slowly - Training takes forever - May get stuck before reaching minimum

Typical Values

Common starting points: 0.1, 0.01, 0.001. Often requires experimentation.

5. Gradient Descent Variants

Batch Gradient Descent

Uses **all** training examples to compute each gradient update.

Pros: - Accurate gradient estimate - Smooth, stable convergence - Guaranteed to converge (for convex problems)

Cons: - Very slow for large datasets - Must fit all data in memory - One update per pass through data

Stochastic Gradient Descent (SGD)

Uses **one** randomly selected example per update.

Pros: - Very fast updates - Can escape local minima (due to noise) - Works with infinite/streaming data

Cons: - Noisy updates, may not converge exactly - High variance in gradient estimates - May oscillate around minimum

Mini-Batch Gradient Descent

Uses a **small batch** of examples (e.g., 32) per update.

Pros: - Balance of speed and stability - Efficient use of GPU parallelism - Most common in practice

Cons: - Introduces batch size as hyperparameter - Still some noise in gradients

Comparison

Method	Samples per Update	Speed	Stability
Batch	All	Slow	High
SGD	1	Fast	Low
Mini-Batch	32-256	Medium	Medium

6. Convergence

When to Stop

Training typically stops when:

- Loss stops decreasing (plateaus)
- Maximum iterations reached
- Validation performance stops improving
- Loss reaches acceptable level

Convergence Criteria

Common approaches:

- Loss change below threshold
- Gradient magnitude below threshold
- Fixed number of epochs

Local vs Global Minima

For convex problems (like linear regression with MSE), there's only one minimum—global. For non-convex problems (neural networks), there may be many local minima. Modern deep learning research suggests local minima are usually fine.

7. Parameters vs Hyperparameters

Parameters

Values **learned from data** during training:

- Weights (w)
- Biases (b)

These define the model's predictions.

Hyperparameters

Values **set before training** by the practitioner:

- Learning rate
- Number of iterations
- Batch size
- Model architecture

These control how training proceeds.

8. The Training Loop

Pseudocode

```
# Initialize
weights = random values
bias = 0

for iteration in range(max_iterations):
    # Forward pass
    predictions = X @ weights + bias

    # Compute loss
    loss = mean((predictions - y)^2)
```

```

# Compute gradients
grad_w = (2/n) * X.T @ (predictions - y)
grad_b = (2/n) * sum(predictions - y)

# Update parameters
weights = weights - learning_rate * grad_w
bias = bias - learning_rate * grad_b

```

Key Components

1. **Forward pass:** Compute predictions with current parameters
 2. **Loss computation:** Measure how wrong we are
 3. **Backward pass:** Compute gradients (how to change parameters)
 4. **Update:** Adjust parameters to reduce loss
-

Key Takeaways

1. **Linear regression:** $y = w \cdot x + b$, simple but powerful
 2. **Loss function:** Measures prediction error (use MSE for regression)
 3. **Gradient:** Vector pointing toward steepest increase
 4. **Gradient descent:** Move opposite to gradient to minimize loss
 5. **Learning rate:** Controls step size (critical hyperparameter)
 6. **Batch/SGD/Mini-batch:** Tradeoffs between speed and stability
 7. **Convergence:** When loss stops decreasing meaningfully
-

Connections to Future Modules

- **Module 3:** What if we have many features? □ Multivariate regression
 - **Module 4:** What about classification? □ Logistic regression (same gradient descent)
 - **Module 6:** What about non-linear functions? □ Neural networks
 - **Module 7:** What about non-gradient methods? □ Decision trees
-

Further Reading

- [Links to foundational papers will be added]
- Sebastian Ruder: “An Overview of Gradient Descent Optimization Algorithms”
- Deep Learning Book, Chapter 4: Numerical Computation